

High Throughput Docking on Helium Cluster

At University of Iowa

Ashley Spies Group

Kenneth TUSSEY

July 23, 2012

1 Getting Started on Helium

Helium is the University of Iowa scientific computing cluster. There are a total of 3748 functional computing cores as of July 23, 2012. Groups are assigned a guaranteed number of cores based on their contribution to the Helium project. Ashley Spies group has 108 computing cores as of July 23, 2012 as part of the *AS* queue. Those 108 are ours to use whenever we please - however, when not in use, they become available as part of the *all.q*, where anyone can use them. If a job happens to be running on one of our processors, and we submit a job to our queue, the offending process will be terminated, allowing us to run our job.

To get started with Helium, you need a Helium account. You can request one at <http://hpc.uiowa.edu/helium-access-agreement>.

Once you have an account, you can login to Helium via ssh.

```
$ ssh -CX username@helium.hpc.uiowa.edu
```

The -CX option allows you to modify text files using gedit. Now that you are logged in to Helium, there are a few important things to know about the environment, backups, etc.

1.1 Modules

Modules are used to change environment variable settings. Modules are loaded into the environment by

```
$ module load modulename
```

where modulename is swapped with whatever module you want to load.

So, if you want to load the Intel libraries into the environment, you would type

```
$ module load intel
```

Common modules in Helium	
openmpi_gnu	Sets MPI variables to point to GNU MPI implementation
openmpi_intel	Sets MPI variables to Intel MPI implementation
intel	Loads variables specific to Intel libraries
python26	Uses Python 2.6.4 interpreter and libraries
python27	Uses Python 2.7 interpreter and libraries

Table 1: Some common helium modules.

If you want to see what modules are loaded, you would use

```
$ module list
```

and if you want to see what modules are available, you would use

```
$ module avail
```

More information is located in the Helium documentation. The good news is that you shouldn't have to worry about the modules to perform the high-throughput docking, but it is good to know how to use modules in case you ever want to tweak something.

1.2 Backups

All home directories are 'backed up' in a sense. Snapshots are taken at specific intervals and saved in a hidden directory in the user's home directory. To get access to your snapshots, you must type

```
$ cd ~/.zfs
```

This will take you into the .zfs directory. Note that if you list the home directory, it will not appear. From the .zfs directory, you will be able to look at old snapshots to recover your data if need be.

1.3 Shared Group Directory

Our group (AS) has a shared directory that all of us can access and read/write files to. It is located in /Groups/AS/. The first thing you want to do when you login to your helium account is create a link in your home folder to this directory.

```
$ ln -s /Groups/AS ~/shared
```

You may call it whatever else you like, but take care that you will need to modify some rc files in order make everything work. After you link the shared directory to your home folder, copy over some rc files into your home directory.

```
$ cp ~/shared/env/*.~
```

This copies a .bashrc and a .chemenv.bashrc file into your home directory, which sets up your user environment variables specific to high-throughput docking. Check to make sure that these are copied over by

```
$ ls -a
```

You should see both files in your home directory. Go ahead and source the .bashrc file.

```
$ source ~/.bashrc
```

2 Setting Up Docking Directory Structure

Once you have the Helium environment setup, we can prepare the docking directory structure. Please note that I have gone to great lengths to minimize the amount of work needed to setup a docking. This will become evident as you read on.

2.1 Ligand Directory

First, start out by making a directory which will contain all of your projects. I like ‘work’, but you may use what you like. For this tutorial, I will use work.

```
$ mkdir ~/work && cd ~/work
```

Now, make a tutorial directory.

```
$ mkdir tutorial && cd tutorial
```

Now, create a directory for the ligand archives. The ligand archives are gzipped archives containing all the individual ligand pdbqt files for a given database.

```
$ mkdir ligand-archives
```

Please name this directory verbatim, or else you will get problems with the docking scripts. The script directly references the ‘ligand-archives’ directory. Now, you will pick a database to screen. The databases that are available are located in ~/shared/ligands/by-vendor. For this tutorial, we will use the analyticon database, because it is small and shouldn’t take very long to perform the entire run. So link or copy (link is preferable for large databases) the archives into the ligand-archives directory.

```
$ ln -s ~/shared/ligands/by-vendor/analyticon/*.tar.gz ligand-archives/
```

Check the ligand-archives folder to make sure you successfully linked the ligand archives.

```
$ ls ligand-archives
```

The links will show up in red if they are broken, and blue if they are competent. This is all that you have to do to setup the ligands part of the directory.

2.2 Receptor Directory

Later, I will write into the tutorial how to create the necessary receptor files using Autogrid, but for the purpose of this tutorial, I will use a pregenerated set of receptor files. In this case, it will correspond to the active site of B. Subtilus Glutamate Racemase (PDB: 1zuw).

Make a directory to contain all of the receptor files.

```
$ mkdir receptor
```

Now, link all of the sample receptor files into the receptor directory.

```
$ ln -s ~/shared/receptor/* receptor/
```

3 Running the Docking

There are three main scripts that run all of the docking. They are called ‘run_parallel-dpf4.sh’, ‘run_parallel-autodock4.sh’, and ‘convert_docking_results.sh’, and are run in that order. The appendix contains the script code for reference. These scripts are located in ~/shared/usr/bin.

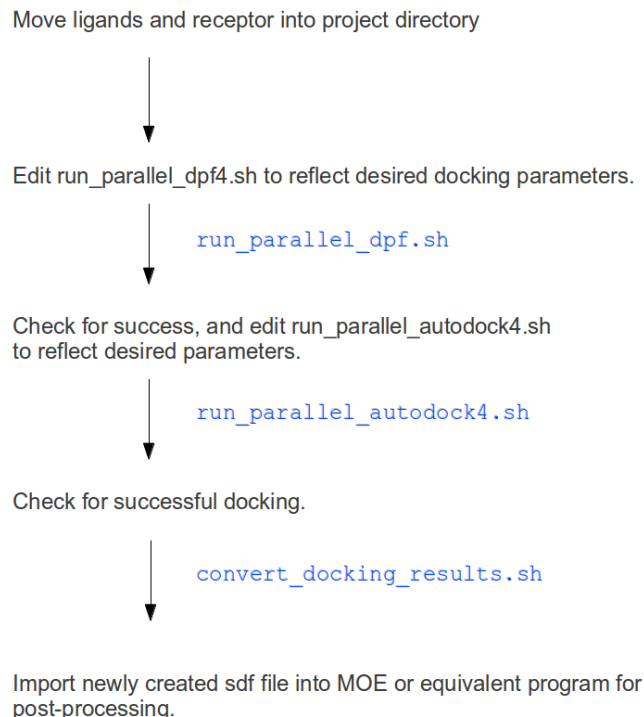


Figure 1: Flowchart for docking

3.1 run_parallel_dp4.sh

This script unpacks the pdbqt archives that you linked to the ligand-archives directory, creates a special folder for each ligand, link the receptor and copy the ligand into the folder, and then create a dpf file that autodock will use perform the docking. Once each folder is created, it is then packed back up into a .tar.gz archive consisting of multiple ligand folders. This may seem counterproductive, since we have to unpack it again to perform the docking, but it is critical to compartmentalize the ligands like this to avoid directory overflow (when we start getting 1000s of folders containing 30 files each, the read/write access becomes painfully slow). We take the hit in overhead, but we will recover it because this script will do this entire procedure in parallel.

In the beginning of this file, you will see about 9 variables. These variables can be changed before running the script. It is important to think about what numbers you want for 'ga_num_evals', 'ga_pop_size', and 'ga_runs'. These are the genetic algorithm controls, and they will greatly influence the accuracy and speed of the docking process. For the purpose of this tutorial, changing it is not necessary. However, for a production screen, you may want to decrease the 'ga_num_evals' to around 100K, and 'ga_num_runs' from 10 to less than 5. They can be changed by editing the file run_parallel_dp4.sh. For the purposes of this tutorial, we should change these settings to as low as possible for the sake of brevity. Change the variables in run_parallel_dp4.sh (this file is in ~/shared/usr/bin - use a text editor to change these) to match Table 2.

Altogether, this script should take around 10 minutes to run for the 22,000 compounds in the analyticon

Settings for docking tutorial	
ga_num_evals	50,000
ga_runs	1

Table 2: Docking parameters for tutorial.

library. For every 108,000 compounds, this script takes around 10 minutes to run (on 108 processors). The script will divide the 22K compounds into 1000 compound segments, and perform the dpf4 creation and archiving in parallel on 23 processors. This creates 23 folders in the ‘ready-for-docking’ folder that is newly created, with around 1K compounds in each.

So for every 1M compounds to process, expect about 1.7 hr to run this phase.

```
$ vim ~/shared/usr/bin/run_parallel.dpf4.sh
```

You may use whatever editor you want (nano is a good choice). Also, take care of anything else that is running on our queue at the time. Check the queue with

```
$ qstat -g c
```

If all 108 cores are not available, then change the variable ‘number_of_processors_anticipated’ to how ever many are available. Please refer to the Appendix Section 1 to see an example of this script.

Now, run the script.

```
$ run_parallel.dpf4.sh
```

When you are finished, you’re directory structure should now look something like this.

```
[bash-ktussey@helium-login-0-0 testrun8]$ ls
dockings  dpf4tmp          ligand-archives  ligands          receptor
dpf4s    ligand_archive_list-3526.txt  ligand-list-for-dpf4  ready-for-docking
```

Figure 2: Snapshot of directory

You should also check the ready-for-docking directory, and make sure that there are 23 files in there.

3.2 run_parallel_autodock4.sh

The next step is to run the script that unpacks the archives in the ‘ready-for-docking’ directory into another directory, and calls the parallel autodock program to dock those compounds. This script also takes the docking results, and puts the dlgs file (the file containing the docked poses and docking specifics) into an archive.

You can see this script in the appendix. There are 3 parameters that can be changed in the beginning of this script, the most important one being how many processors you would like to use. Again, check the AS queue to see how many you have, and run on all 108 if it is possible. If all 108 are available, than you do not need to modify this script.

Now, run the parallel docking.

```
$ run_parallel_autodock4.sh
```

Now, this script will take a while to finish. This time is very sensitive to the nature of the library being docked. Libraries with compounds containing many rotatable bonds will take noticeably longer to run than libraries with a low number of rotatable bonds. Increasing the number of evaluations and runs will cause the job to run longer. Balancing all of these variables is important to running a succesful, yet timely screen. When this is finished, go into the ‘docking_results’ directory, and check that there are results .tar.gz files. There should be 23 archives. You can check them with the following script.

```

1 #!/bin/bash
  num_results=0
3 for i in docking_results
  do
5     num_results=$((num_results+$(zcat $i | tar -tf - | wc -l)))
  done
7 echo $num_results

```

It should be around 22K, depending on how many ligands autodock choked on. In this tutorial, I successfully docked 22,446 compounds in 72 minutes. 27 minutes were spent in overhead, and 45 minutes in docking. You can check these timings in the file timings.txt. The overhead time seems high for this tutorial, but when we do production runs using more intensive docking parameters, this will be much smaller proportionally.

3.3 convert_docking_results.sh

Now that autodock has finished, you have a nice set of docked structures. Unfortunately, they are locked up in a dlg file, which no commercial or academic program (to my knowledge) has taken the time to include in their supported formats. Luckily, with a few scripting tricks, we can extract the relevant pdbqt files from the dlg file, and then convert them to sdf using a custom program that relies on OpenBabel C++ bindings. To convert the docking results into an SDF format, run

```
$ convert_docking_results.sh
```

This script currently processes the results serially. However, it is trivial to parallelize. If this becomes too cumbersome to run, it can be parallelized to 100x speedup. For 22K compounds, it takes around 10 minutes to run.

When this script is finished, you can find the docking results in docking_results/final_results.sdf.

You can take this file, and import into a MOE database, or use OpenEye's vina to view the results. The receptor is found in receptor/.

4 Appendix

4.1 run_parallel_dpf4.sh

```
#!/bin/bash
2

4 # Some control numbers for the script , does not need to be changed. Can be if
  # wanted, for optimization purposes
6
  sge_name_for_script=prepdpf4
8 number_of_processors_anticipated=108
  max_ligand_list_size=2000
10 min_ligand_list_size=1000
  limit_num_ligands=100000
12 ga_num_evals=500000
  ga_pop_size=150
14 ga_runs=10
  sge_queue_custom=AS
16

18 T="$(date +%s)"
  # Need to know how many ligands we will be dealing with
20
  let global_num_ligands=0
22 let global_total_num_ligands=0

24 if [ $((($(find ligand-archives -name "*.tar.gz" | wc -l))) -gt 0 ) ]; then sleep 0
  else
26   echo "No ligand archives found. Exiting"
    exit
28 fi

30

32 echo "Getting total ligand count..."
  for i in ligand-archives/*.tar.gz
34 do

36   global_num_ligands=$(zcat $i | tar -tf - | wc -l)
    global_total_num_ligands=$((global_total_num_ligands+global_num_ligands))
38 done
  echo "Total ligand count: $global_total_num_ligands"
40

42

44 # Function that creates a list of ligand files , divied up by how many processors and compounds
  there are.

46 function prepare_for_dpf4 () {

48   if [ $((total_num_ligands/number_of_processors_anticipated)) -gt $max_ligand_list_size ]; then
    let ligand_list_size=$max_ligand_list_size
49   elif [ $((total_num_ligands/number_of_processors_anticipated)) -lt $min_ligand_list_size ]; then
    let ligand_list_size=$min_ligand_list_size
50   else
    let ligand_list_size=$((total_num_ligands/$number_of_processors_anticipated+1))
52   fi
54

56   cd ligands
    let counter_l=0
58   let counter_j=1

60   if [ -d ../ligand-list-for-dpf4 ]; then rm ../ligand-list-for-dpf4/*
  else
62   mkdir ../ligand-list-for-dpf4
```

```

64     fi
66 # put each ligand file's name into a list of ~2K compounds,
67 # depending on number of procs and number of ligands
68 for i in *.pdbqt
69 do
70     echo $i >> ../ligand-list-for-dpf4/ligand-list-for-dpf4-${counter_l}
71     if [ $((counter_j % $ligand_list_size)) -eq 0 ]; then
72         counter_l=$((counter_l+1))
73     fi
74     counter_j=$((counter_j+1))
75 done
76 cd ../
77 }
78
80
82 # This function creates a shell script that will be submitted to SGE to process the ligands into
83 # folders which are meaningful to the parallel autodock. For each list of ~2K ligands produced in
84 # the function prepare_for_dpf4, a shell script is spawned that will iterate through the list. Each
85 # of these shell scripts is submitted to whichever queue is specified in the control portion of
86 # this script.
87
88 function run_dpf4 () {
89
90     if [ -d dpf4tmp ]; then rm dpf4tmp/*
91     else
92         mkdir dpf4tmp
93     fi
94
95     basename=$PWD
96
97     if [ ! -d ready-for-docking ]; then mkdir ready-for-docking; fi
98     if [ ! -d Dockings ]; then mkdir Dockings; fi
99     if [ ! -d dpf4s ]; then mkdir dpf4s; fi
100
101     cd ligand-list-for-dpf4
102     for i in ligand-list-for-dpf4-*
103     do
104         touch ../dpf4tmp/$i.sh
105         echo "#!/bin/bash" >> ../dpf4tmp/$i.sh
106         echo "module load intel" >> ../dpf4tmp/$i.sh
107         echo "module load python26" >> ../dpf4tmp/$i.sh
108         echo 'random_folder=$RANDOM' >> ../dpf4tmp/$i.sh
109         echo 'mkdir Dockings/docking-${random_folder}' >> ../dpf4tmp/$i.sh
110         echo -ne 'for j in $(cat ' >> ../dpf4tmp/$i.sh
111         echo -ne "ligand-list-for-dpf4/$i" >> ../dpf4tmp/$i.sh
112         echo ');' >> ../dpf4tmp/$i.sh
113         echo 'do ' >> ../dpf4tmp/$i.sh
114         echo '    ligand_name=$(basename $j .pdbqt)' >> ../dpf4tmp/$i.sh
115         echo '    receptor_name=$(basename receptor/*.pdbqt .pdbqt)' >> ../dpf4tmp/$i.sh
116         echo -ne '    prepare_dpf42.py -l ligands/${j} -r receptor/1zuw.pdbqt -o dpf4s/${ligand_name}-${receptor_name}.dpf' >> ../dpf4tmp/$i.sh
117         echo "    -p ga_num_evals=$ga_num_evals -p ga_pop_size=$ga_pop_size -p ga_run=$ga_runs -p rmstol=2.0" >> ../dpf4tmp/$i.sh
118         echo '    mkdir Dockings/docking-${random_folder}/${ligand_name}-${receptor_name}' >> ../dpf4tmp/$i.sh
119         echo '    mv dpf4s/${ligand_name}-${receptor_name}.dpf Dockings/docking-${random_folder}/${ligand_name}-${receptor_name}' >> ../dpf4tmp/$i.sh
120         echo -ne "    ln -s $basename/" >> ../dpf4tmp/$i.sh
121         echo 'receptor/* Dockings/docking-${random_folder}/${ligand_name}-${receptor_name}' >> ../dpf4tmp/$i.sh
122         echo '    mv ligands/${ligand_name}.pdbqt Dockings/docking-${random_folder}/${ligand_name}-${receptor_name}' >> ../dpf4tmp/$i.sh

```



```

124     echo 'done' >> ../dpf4tmp/$i.sh
125     echo 'cd Dockings/docking-${random_folder}' >> ../dpf4tmp/$i.sh
126     echo 'tar -czf ../../ready-for-docking/docking_folder-${random_folder}.tar.gz * --remove-files '
127     >> ../dpf4tmp/$i.sh
128     echo 'cd ../../' >> ../dpf4tmp/$i.sh
129     echo 'rm -rf Dockings/docking-${random_folder}' >> ../dpf4tmp/$i.sh
130
131 done
132
133 cd ../
134
135 # Submit each shell script to the cluster as an independent, serial job. When the lustre file
136 # system becomes
137 # available again, we will change the above script to output to lustre (or create a symlink from
138 # the folder
139 # dpf4s to some folder on lustre.
140 for i in dpf4tmp/*.sh
141 do
142     qsub -V -N $sge_name_for_script -cwd -j y -o /nfsscratch/$sge_name_for_script.log -q
143     $sge_queue_custom $i
144 done
145
146 # This line queries the queue, looks for the job named $sge_name_for_script,
147 # and looks every 20 seconds until there are no
148 # more jobs with then name $sge_name_for_script. The function then exits.
149 echo "Running jobs on cluster..."
150 sleep 10
151 job_waiter=$(qstat | grep $sge_name_for_script)
152 while [ "$job_waiter" ]; do sleep 20; job_waiter=$(qstat | grep $sge_name_for_script); done
153
154 echo "Finished!"
155 }
156
157
158
159
160 let number_of_ligands=0
161 let total_number_of_ligands=0
162 let total_docking_archives=1
163 let ligand_counter_for_output=0
164
165 ligand_archive_list_notifier=$RANDOM
166
167
168 if [ -d ligands ]; then
169     rm -rf ligands/*
170 else
171     mkdir ligands
172 fi
173
174
175
176 for i in ligand-archives/*.tar.gz
177 do
178     echo $i >> ligand_archive_list-${ligand_archive_list_notifier}.txt
179 done
180
181 # Typically, each archive will house 100K compounds in pdbqt format. Typically, we have around 50
182 # processors
183 # free at any given time. This means that the best way to divy up each archive is to split it into
184 # subsets
185 # containing ~2K ligands apiece. The time it takes to process 2K compounds is around 10 minutes.
186
187 exec 3<&0

```

```

186 exec 0<ligand_archive_list-${ligand_archive_list_notifier}.txt
188 while read line
189 do
190
192 # We will fill our ligand directory up to 100K ligands, but no more so that file access isn't
    slowed down
    if [ $total_number_of_ligands -lt $limit_num_ligands ]; then
194         echo "Extracting $line..."
        tar -xzf $line -C ligands
196         number_of_ligands=$(zcat $line | tar -tf - | wc -l)
        ligand_counter_for_output=$((ligand_counter_for_output+number_of_ligands))
198         total_number_of_ligands=$((total_number_of_ligands+number_of_ligands))
        #sleep 0
200     else
        echo "Preparing dpf4s for $ligand_counter_for_output out of $global_total_num_ligands ligands"
202         prepare_for_dpf4
        run_dpf4
204         echo "Extracting $line..."
        tar -xzf $line -C ligands
206         number_of_ligands=$(zcat $line | tar -tf - | wc -l)
        ligand_counter_for_output=$((ligand_counter_for_output+number_of_ligands))
208         total_number_of_ligands=$number_of_ligands
        total_docking_archives=$((total_docking_archives+1))
210     fi
212
213 done
214
216 # Finally, perform the last run through of ligands for processing.
    prepare_for_dpf4
218 run_dpf4
220
221 exec 0<&3
222
223 #echo "ligands/${LIGAND_FILENAME}"
224 #echo "dpf4s/${basename $LIGAND_FILENAME}.pdbqt"
226
227 if [ ! -d logs ]; then mkdir logs; fi
228 rm -rf dpf4tmp
    rm -rf dpf4s
230 rm -rf ligand_archive_list-${ligand_archive_list_notifier}.txt
    rm -rf ligand-list-for-dpf4
232
233 T="$(date +%s)-T)"
234
235 echo "Processed $global_total_num_ligands ligands" > timings.txt
236 echo "Timing in minutes to process: $(T/60)" >> timings.txt

```

4.2 run_parallel_autodock4.sh

```
#!/bin/bash
2

4 num_processors_for_docking=108
  sge_name_for_script=ad42_mpi
6 sge_name_for_queue=AS

8 let overhead_time=0
  let docking_time=0
10 let total_ligands=0
  let log_indicator=$RANDOM
12 let counter=0
  ls ready-for-docking > undocked_list.txt
14 let total_files=$(cat undocked_list.txt | wc -l)

16 for i in $(cat undocked_list.txt)
do
18
    total_ligands=$((total_ligands + $(zcat ready-for-docking/$i | tar -tf - | grep dpf | wc -l)))
20 echo "Processing ligands up to $total_ligands compounds"

22 echo "Processing $counter of $total_files"
  counter=$((counter+1))
24 T=$(date +%s)
  basename=$(basename $i .tar.gz)
26 if [ ! -d Dockings ]; then mkdir Dockings; fi
  if [ ! -d docking-results ]; then mkdir docking-results; fi
28 tar -xzf ready-for-docking/$i -C Dockings

30 ls Dockings/ > ligand_list.txt

32 echo '#!/bin/bash' > cluster_sub_script.sh
  echo "module load intel" >> cluster_sub_script.sh
34 echo "module load openmpi-intel" >> cluster_sub_script.sh
  echo 'mpirun autodock-intel-serial-mpi $PWD/ligand_list.txt $PWD/Dockings/ $PWD/logs/
    default_seed reuse_maps' >> cluster_sub_script.sh

36 T=$((($(date +%s)-T))
38 overhead_time=$((overhead_time+T))

40 T=$(date +%s)
  qsub -V -N $sge_name_for_script -pe orte $num_processors_for_docking -cwd -j y -o /nfsscratch/
    autodocklogs-{$log_indicator}.log -q $sge_name_for_queue cluster_sub_script.sh
42

44 echo "Running jobs on cluster..."
  job_waiter=$(qstat | grep $sge_name_for_script)
46 while [ "$job_waiter" ]; do sleep 20; job_waiter=$(qstat | grep $sge_name_for_script); done

48 echo "Finished!"
  T=$((($(date +%s)-T))
50 docking_time=$((docking_time+T))

52 T=$(date +%s)
54 echo "Packing docking logs..."
  mkdir tmp_dlglinks

56
  for j in $(find $PWD/Dockings -name "*.dlg")
  do
58      ln -s $j tmp_dlglinks
60  done

62 cd tmp_dlglinks
  tar -czhf ${basename}-results.tar.gz * --remove-files
64 mv ${basename}-results.tar.gz ../docking_results
```

```

66  cd ../
    rm -rf tmp_dlglinks

68  echo "Finished!"

70  echo "Cleaning up docking directory..."

72  rm -rf Dockings/*

74  echo "Finished!"

76  T="$(date +%s)-T)"
    overhead_time="$(overhead_time+T)"

78  echo $i >> finished_docking_list.txt
80  sed -i 'ld' undocked_list.txt

82  done

84  total_time="$(docking_time+overhead_time)"

86  echo "Docked $total_ligands ligands in a total of $((total_time/60)) minutes" >> timings.txt

88  echo "Overhead time = $((overhead_time/60)) minutes" >> timings.txt
    echo "Docking time = $((docking_time/60)) minutes" >> timings.txt

```

4.3 convert_docking_results.sh

```
2
3  #!/bin/bash
4
5
6  counter=1;
7  total=$(ls docking_results | wc -l)
8  for i in docking_results/*.tar.gz
9  do
10
11
12     if [ ! -d docking_results/tmp ]
13     then mkdir docking_results/tmp
14     else
15         rm docking_results/tmp/*
16         echo "Cleaning up..."
17     fi
18
19
20
21
22     echo "Converting $counter of $total docking results..."
23
24     echo "Untarring docking results..."
25     tar -xzf $i -C docking_results/tmp
26
27     echo "Converting dlgs to sdfs..."
28     dlgtosdf -i docking_results/tmp/*.dlg >> docking_results/final_results.sdf
29
30     mv $i processed_results
31
32     counter=$((counter+1))
33 done
```